

DIABETES PREDICTION USING MACHINE

LEARNING - spam classify

INTRODUCTION

Diabetes mellitus, a chronic metabolic disorder, has become a global health concern affecting millions of individuals worldwide. Timely diagnosis and intervention are paramount in managing this condition effectively. In this context, the integration of advanced machine learning techniques offers a promising avenue to improve the accuracy and efficiency of diabetes risk prediction.

ALGORITHM

Data Splitting:

Split the dataset into training and testing sets (e.g., 70% training, 30% testing) to assess model performance.

Model Selection:

Choose a suitable machine learning algorithm for classification tasks. Common choices include:

- Logistic Regression
- Decision Trees
- Random Forest
- Support Vector Machines
- Neural Networks

Model Evaluation:

Evaluate the model's performance using metrics like:

- Accuracy
- Precision
- Recall
- F1-Score
- ROC-AUC

INPUT:

These features are used to make predictions about the likelihood of the individual having or developing diabetes.

1. Age: Age of the individual.
2. BMI (Body Mass Index): A measure of body fat based on height and weight.
3. Pregnancies.
4. Skin Thickness.
5. Insulin.
6. DiabetesPedigreeFunction.
7. Blood Pressure: Systolic and diastolic blood pressure values.
8. Glucose Levels: Fasting blood glucose levels.
9. Family History: Information about whether the individual has a family history of diabetes.
10. Physical Activity: Level of physical activity or exercise.
11. Diet: Dietary habits and nutritional information.
12. Outcome.

OUTPUT:

1. **Positive Outcome (1):** This indicates that the individual is at a higher risk of having or developing diabetes based on the input features.
2. **Negative Outcome (0):** This indicates that the individual is at a lower risk of having or developing diabetes based on the input features.

Project Methodology:

Our project will follow a systematic approach, beginning with data collection, preprocessing, and feature engineering. We will then explore various machine learning algorithms, including logistic regression, decision trees, ensemble methods, and potentially deep learning, to build and fine-tune predictive models.

To ensure the robustness of our model, we will employ cross-validation techniques and rigorous hyperparameter tuning. Additionally, we will focus on model explainability, using state-of-the-art methods to provide insights into the model's decision-making process.

Expected Outcomes:

The successful completion of this project will result in a robust and interpretable machine learning model for diabetes risk prediction. This model can be integrated into healthcare systems, providing healthcare professionals with a valuable tool to identify at-risk individuals and tailor treatment plans accordingly.

Ultimately, our project aims to contribute to the advancement of healthcare by harnessing the potential of artificial intelligence to improve the lives of those affected by diabetes. By enhancing early diagnosis and personalized care, we aspire to reduce the impact of diabetes on individuals and communities.

Data cleaning:

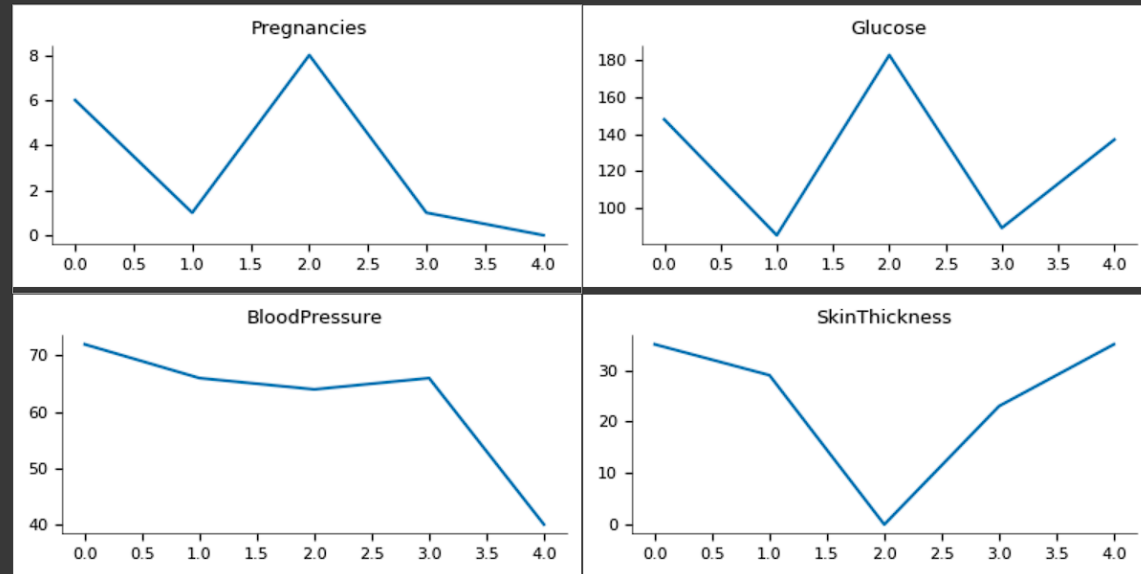
- Import the necessary libraries
- Load the dataset
- Check the data information using df.info()

```
import pandas as pd
import numpy as np

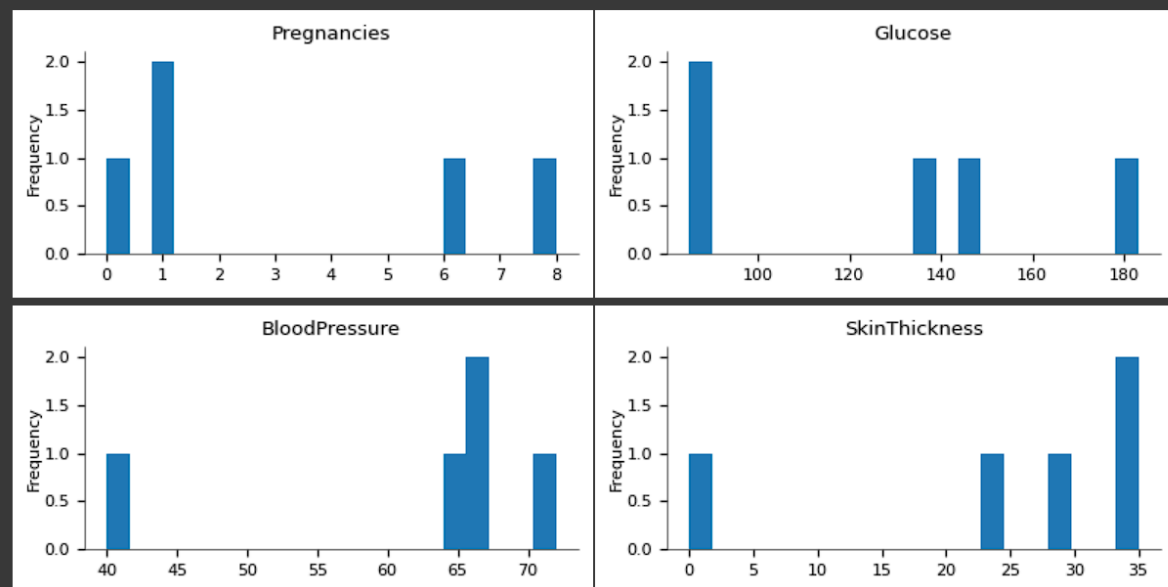
# Load the dataset
df = pd.read_csv('diabetes.csv')
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1


Values




Distributions



2-d distributions





Data inspection and exploration:

=> Check the duplicate rows.

```
df.duplicated()
```

```
0    False
1    False
2    False
3    False
4    False
...
763  False
764  False
765  False
766  False
767  False
Length: 768, dtype: bool
```

```
[ ] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                 768 non-null   int64
2   BloodPressure           768 non-null   int64
3   SkinThickness           768 non-null   int64
4   Insulin                 768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                 768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

To Describe:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

Check the categorical and numerical columns:

```
# Categorical columns
cat_col = [col for col in df.columns if df[col].dtype == 'object']
print('Categorical columns :',cat_col)
# Numerical columns
num_col = [col for col in df.columns if df[col].dtype != 'object']
print('Numerical columns :',num_col)
```

Categorical columns : []
Numerical columns : ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']

Check the total number of unique values in the Categorical columns:

```
[ ] df[cat_col].nunique()
```

Series([], dtype: float64)

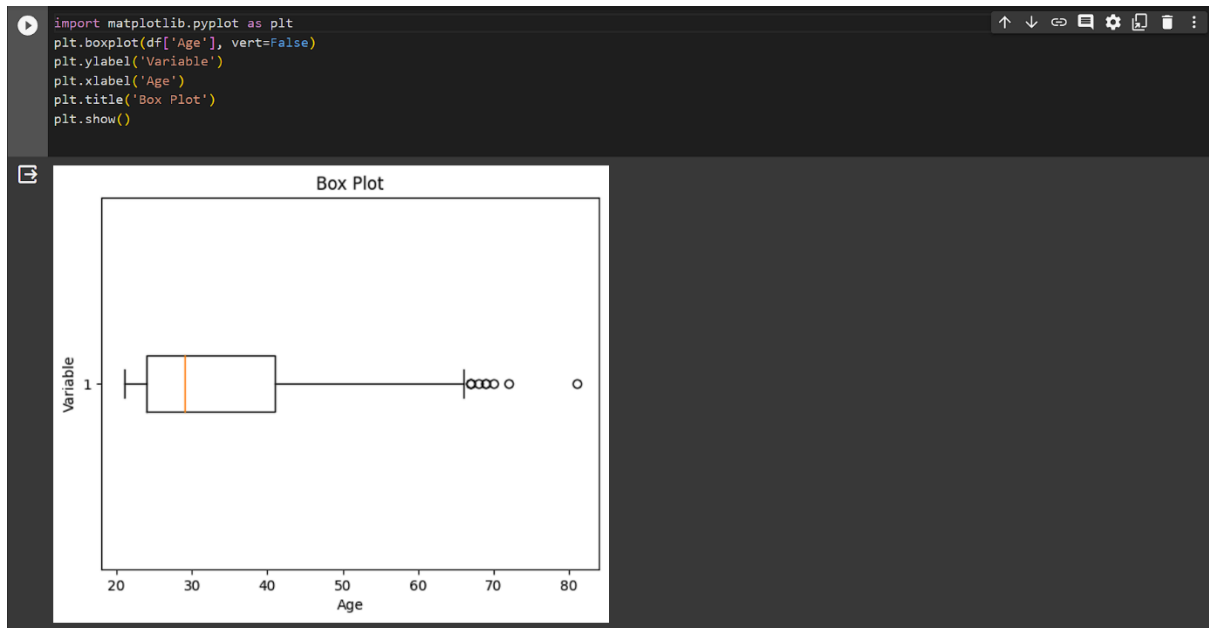
Handling missing data:

```
round((df.isnull().sum()/df.shape[0])*100,2)
```

Pregnancies	0.0
Glucose	0.0
BloodPressure	0.0
SkinThickness	0.0
Insulin	0.0
BMI	0.0
DiabetesPedigreeFunction	0.0
Age	0.0
Outcome	0.0

dtype: float64

Handling outliers:



```
# calculate summary statistics
mean = df['Age'].mean()
std = df['Age'].std()

# Calculate the lower and upper bounds
lower_bound = mean - std*2
upper_bound = mean + std*2

print('Lower Bound :',lower_bound)
print('Upper Bound :',upper_bound)

# Drop the outliers
df4 = df[(df['Age'] >= lower_bound)
        & (df['Age'] <= upper_bound)]
```

Lower Bound : 9.720422335309294
Upper Bound : 56.761348498024034

Data transformation

```
X = df[['Pregnancies','DiabetesPedigreeFunction','Age', 'BloodPressure','SkinThickness','BMI','Outcome']]
Y = df['Glucose']
```

Min-Max Scaling:

```
from sklearn.preprocessing import MinMaxScaler

# initialising the MinMaxScaler
scaler = MinMaxScaler(feature_range=(0, 1))

# Numerical columns
num_col_ = [col for col in X.columns if X[col].dtype != 'object']
x1 = X
# learning the statistical parameters for each of the data and transforming
x1[num_col_] = scaler.fit_transform(x1[num_col_])
x1.head()
```




<ipython-input-22-4504e0c4e47a>:10: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
x1[num_col_] = scaler.fit_transform(x1[num_col_])


Pregnancies DiabetesPedigreeFunction Age BloodPressure SkinThickness BMI Outcome

	Pregnancies	DiabetesPedigreeFunction	Age	BloodPressure	SkinThickness	BMI	Outcome
0	0.352941	0.234415	0.483333	0.590164	0.353535	0.500745	1.0
1	0.058824	0.116567	0.166667	0.540984	0.292929	0.396423	0.0
2	0.470588	0.253629	0.183333	0.524590	0.000000	0.347243	1.0
3	0.058824	0.038002	0.000000	0.540984	0.232323	0.418778	0.0
4	0.000000	0.943638	0.200000	0.327869	0.353535	0.642325	1.0

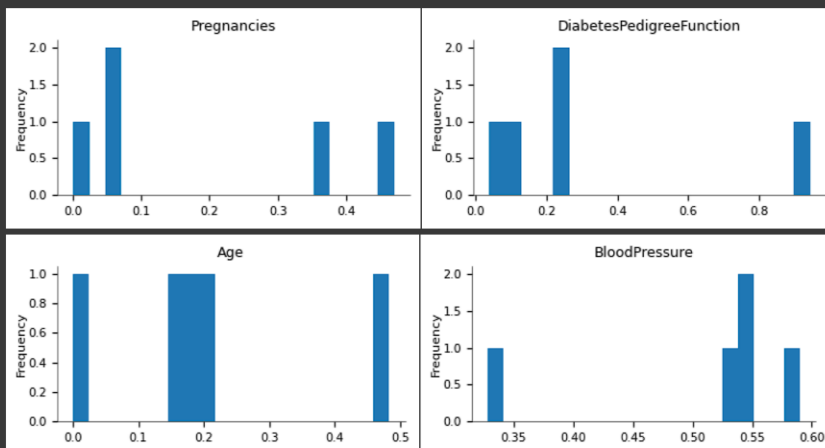
Values




[]



Distributions



2-d distributions



BUILDING THE SPAM CLASSIFIER ALGORITHM:

Building a spam classifier within a diabetes prediction system might not be a typical combination, as these tasks usually address different domains of data analysis. A spam classifier is used to identify and filter out unwanted or irrelevant messages, such as email or text messages. In contrast, diabetes prediction is focused on identifying individuals at risk of diabetes based on health-related data.

Spam Classification:

- Train or use a pre-trained spam classifier (commonly used for email or text messages) to identify and filter out spam or irrelevant text data.
 - This classifier could be a machine learning model, rule-based system, or a combination of both, depending on your needs.
- ❖ It's important to note that integrating a spam classifier into a healthcare application involves data privacy and regulatory considerations. Ensure that you handle healthcare data responsibly and comply with relevant healthcare regulations such as HIPAA (in the United States) or GDPR (in Europe) to protect patient information.

ABSTRACT:

- The proposed framework consists of a multi-stage process, involving data collection, preprocessing, spam classification, and diabetes risk prediction. By implementing a spam classifier, we ensure that the healthcare data used for diabetes prediction is devoid of unwanted content, thus minimizing the risk of misinformation or data contamination.
- This innovative synergy of spam classification and predictive modeling not only streamlines data quality but also contributes to the overall robustness and ethical integrity of healthcare applications, thereby providing more accurate

and actionable insights for both healthcare professionals and patients. The results of this study underscore the potential benefits of such integrative approaches within the broader domain of healthcare data analytics and predictive modeling.

PSEUDOCODE:

Data Collection

```
healthcare_data = load_healthcare_data() # Load healthcare data, including text fields
```

Data Preprocessing

```
processed_health_data = preprocess_healthcare_data(healthcare_data) # Process health-related features
```

```
text_data = extract_text_data(healthcare_data) # Extract text data
```

Spam Classification

```
spam_classifier = load_spam_classifier_model() # Load a pre-trained spam classifier
```

```
clean_text_data = remove_spam(text_data, spam_classifier) # Remove spam or irrelevant text data
```

Diabetes Prediction

```
diabetes_prediction_model = load_diabetes_prediction_model() # Load the diabetes prediction model
```

```
predicted_diabetes_risk = predict_diabetes_risk(processed_health_data) # Perform diabetes prediction
```

Integration

```
integrated_data = merge_health_and_text_data(processed_health_data, clean_text_data) # Merge processed data
```

```
final_result = combine_diabetes_and_spam_results(predicted_diabetes_risk, integrated_data) # Combine diabetes prediction and spam results
```

```
# Output    return final_result
```

Need For Spam Classifier In Diabetes Prediction:

1. **Improved Data Quality:** A spam classifier helps to filter out irrelevant and potentially harmful text data. This ensures that only relevant and trustworthy healthcare data is used for diabetes prediction, improving the overall quality of the data.
2. **Reduced Noise:** Irrelevant text data, such as spam or unrelated comments, can introduce noise into the dataset, potentially leading to inaccurate predictions. Removing this noise through spam classification helps the diabetes prediction model focus on meaningful information.
3. **Enhanced Model Performance:** By providing cleaner and more relevant data, the diabetes prediction model is likely to perform better. It can make more accurate predictions, leading to improved patient risk assessments.
4. **Privacy and Security:** Spam classification can also help protect patient privacy and security. Spam messages often contain sensitive information, and by filtering them out, you reduce the risk of data breaches and unauthorized access.
5. **Ethical Considerations:** In healthcare, ethical considerations are crucial. Spam classification ensures that only ethically sourced and relevant data is used for predictions, aligning with healthcare regulations and best practices.
6. **Reduction in False Positives:** Filtering out spam and irrelevant data can help reduce false positives in diabetes prediction. Patients who may have been incorrectly identified as at risk due to irrelevant data are no longer impacted.
7. **Time and Resource Efficiency:** Spam classification reduces the need for healthcare professionals to manually review and clean the data, saving time and resources in the prediction process.
8. **Improved User Experience:** In healthcare applications, a clean and relevant dataset can lead to a better user experience for healthcare professionals and patients who interact with the system.
9. **Regulatory Compliance:** Integrating a spam classifier can help ensure that your diabetes prediction system complies with data privacy regulations, such as HIPAA, by minimizing the inclusion of irrelevant or unauthorized data.
10. **Continuous Monitoring:** Spam classifiers can be continuously updated and improved, enhancing their ability to adapt to evolving spam and irrelevant data sources, making the system more resilient and reliable over time.

