

ARTIFICIAL INTELLIGENCE

Project Title: AI-BASED DIABETES PREDICTION SYSTEM

Phase 3:Data visualization

INTRODUCTION:

Data visualization plays a pivotal role in the field of Machine Learning (ML). It involves representing and presenting data in graphical or visual formats to gain insights, understand patterns, and communicate findings effectively.

We will select relevant features that can impact diabetes risk prediction.

Data Visualization Algorithm:

1. Data Preprocessing
2. Data Exploration
3. Plotting
4. Customization
5. Interactivity

DATASET:

Dataset Link: <https://www.kaggle.com/datasets/mathchi/diabetes-data-set>

Algorithm: Data Preprocessing, Classification, and visualisation.

1. Import the necessary libraries (NumPy, pandas, scikit-learn, and Matplotlib).
2. Define a list of relevant features (e.g., 'Glucose', 'BloodPressure', 'BMI', 'Age').
3. Load the diabetes dataset from a CSV file, and subset it to include relevant features and the 'Outcome' target variable.
4. Standardize the feature data using StandardScaler.
5. Split the data into training and testing sets.
6. Create a Support Vector Machine (SVM) classifier with a linear kernel and train it on the training data.
7. Make predictions on both the training and testing sets and calculate accuracy scores.
8. Visualize accuracy using a bar chart.
9. Create a histogram to visualize the distribution of the "Glucose" feature.
10. Make predictions on a sample input data point (e.g., a hypothetical patient's data).

11. Generate synthetic data:
 - a. Determine the number of synthetic samples to generate.
 - b. Initialize lists to store synthetic data points and labels.
 - c. For each synthetic sample, randomly select an index from the real data.
 - d. Select a real data point and its label.
 - e. Create a slightly modified version of the real data point with random noise. f. Append the modified data point and its label to the synthetic data.
12. Combine real and synthetic data.
13. Print the first 5 samples of the synthetic data and their corresponding labels.

PROGRAM:

Import necessary libraries

```
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
```

Define a list of relevant

```
features relevant_features = ['Glucose',
                              'BloodPressure', 'BMI', 'Age']
```

Load the diabetes dataset from a CSV file (Replace 'diabetes.csv' with the actual dataset path)

```
diabetes_dataset = pd.read_csv('diabetes.csv')
```

```
# Subset the dataset to include relevant features and the 'Outcome' target variable diabetes_dataset
= diabetes_dataset[relevant_features + ['Outcome']]
```

Split the data into features (X) and the target (Y)

```
X = diabetes_dataset[relevant_features]
```

```
Y = diabetes_dataset['Outcome']
```

Standardize the feature data using StandardScaler

```
scaler = StandardScaler()
```

```
X = scaler.fit_transform(X)
```

Split the data into training and testing sets

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=2)
```

Create an SVM classifier with a linear kernel

```
classifier=svm.SVC(kernel='linear') classifier.fit(X_train,  
Y_train)
```

Make predictions on the training and testing sets

```
Y_train_pred = classifier.predict(X_train)
```

```
Y_test_pred = classifier.predict(X_test)
```

Calculate accuracy scores for training and testing data

```
training_data_accuracy  
= accuracy_score(Y_train_pred, Y_train) test_data_accuracy =  
accuracy_score(Y_test_pred, Y_test)
```

```
# Visualize accuracy using a bar chart accuracies = [training_data_accuracy,  
test_data_accuracy] labels = ['Training Data', 'Test Data'] plt.figure(figsize=(8, 6))  
plt.bar(labels, accuracies, width=0.4, align='center', alpha=0.5, color=['blue',  
'green']) plt.xlabel('Data Split') plt.ylabel('Accuracy') plt.title('Training and Test  
Accuracies') # Data visualization section: Histogram of the "Glucose" feature
```

```
plt.figure(figsize=(8, 6)) plt.hist(X[:, 0], bins=20, color='blue',  
alpha=0.7) plt.xlabel('Glucose Level') plt.ylabel ('Frequency')  
plt.title('Distribution of Glucose Levels') # Make predictions on a
```

```
sample input data point input_data = np.array([1, 85, 66, 29, 0,  
26.6, 0.351, 31]).reshape(1, 1) input_data =  
scaler.transform(input_data) prediction =  
classifier.predict(input_data)
```

```
# Number of synthetic samples to generate num_samples_to  
generate = 500
```

```
# Initialize lists to store synthetic data synthetic_data  
= [] synthetic_labels =  
[]
```

Generate synthetic data points based on existing data for

_ in range(num_samples_to_generate):

```
    random_index = np.random.randint(0, len(X))    real_data_point = X[random_index]    real_label
= Y[random_index]    modified_data_point = real_data_point + np.random.normal(0,
0.1, size=real_data_point.shape)    synthetic_data.append(modified_data_point)
synthetic_labels.append(real_label)
```

Combine real and synthetic data

```
X_synthetic = np.vstack([X, np.array(synthetic_data)])
```

```
Y_synthetic = np.concatenate([Y, np.array(synthetic_labels)])
```

Print the first 5 samples of the synthetic data

```
print("Synthetic Data (X_synthetic):") print(X_synthetic[:5]) #
```

Print the corresponding labels for the first 5 samples


```
print("Synthetic Labels (Y_synthetic):") print(Y_synthetic[:5])
```

OUTPUT:

```
Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6       148           72           35         0   33.6
1           1        85           66           29         0   26.6
2           8       183           64            0         0   23.3
3           1        89           66           23         94   28.1
4           0       137           40           35       168   43.1
..         ...       ...           ...           ...       ...   ...
763         10       101           76           48       180   32.9
764         2       122           70           27         0   36.8
765         5       121           72           23       112   26.2
766         1       126           60            0         0   30.1
767         1        93           70           31         0   30.4

DiabetesPedigreeFunction  Age
0                0.627    50
1                0.351    31
2                0.672    32
3                0.167    21
4                2.288    33
..               ...     ...
763              0.171    63
764              0.340    27
765              0.245    30
766              0.349    47
767              0.315    23

[768 rows x 8 columns]
0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64
[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
    1.4259954 ]
 -0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
 -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
 -0.10558415]
 ...
 [ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
 -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
  1.17073215]
 [-0.84488505 -0.8730192  0.04624525 ... -0.20212881 -0.47378505
 -0.87137393]]
 [[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
    1.4259954 ]
 -0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
 -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
 -0.10558415]
 ...
 [ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
 -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
  1.17073215]
 [-0.84488505 -0.8730192  0.04624525 ... -0.20212881 -0.47378505
 -0.87137393]]
0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64
(768, 8) (614, 8) (154, 8)
Accuracy score of the training data : 0.7866449511400652
```



```
Accuracy score of the test data : 0.7727272727272727
[[1.00000000e+00 8.50000000e+01 6.60000000e+01 2.90000000e+01
 3.00685403e-17 2.66000000e+01 3.51000000e-01 3.10000000e+01]]
[1]
The person is diabetic
Synthetic Data (X_synthetic):
[[ 0.63994726  0.84832379  0.14964075  0.90726993 -0.69289057  0.20401277
  0.46849198  1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575  0.53090156 -0.69289057 -0.68442195
 -0.36506078 -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 -1.28821221 -0.69289057 -1.10325546
  0.60439732 -0.10558415]
 [-0.84488505 -0.99820778 -0.16054575  0.15453319  0.12330164 -0.49404308
 -0.92076261 -1.04154944]
 [-1.14185152  0.5040552  -1.50468724  0.90726993  0.76583594  1.4097456
  5.4849091  -0.0204964 ]]
Synthetic Labels (Y_synthetic):
[1 0 1 0 1]
```

Result:

The code preprocesses a diabetes dataset, trains a support vector machine (SVM) classifier, visualizes the distribution of glucose levels(visualization), and augments the dataset with synthetic data for classification and analysis. Specific results depend on the dataset and problem domain, including classifier accuracy and improved data quality for enhanced model performance.